

Everything you need to know about ChainRules 1.0

Miha Zgubič, Research Software Engineer @  INVENIA
LABS





We use machine learning to optimise the electricity grids.

more Julia, less emissions

Come join us!



ChainRules has many



Alex Arslan



David Widmann



Jarrett Revels




Lyndon White



Michael Abbott



Miha Zgubic



Matt Brzezinski



Nick Robinson



Seth Axen



Simeon Schaub



Will Tebbutt



Yingbo Ma



...many wonderful contributors


Shashi Gowda
Pietro Vertech
Kristoffer Carlsson
Andrew Fitzgibbon
Curtis Vogt
Jerry Ling
Mason Protter
Steven Johnson
Viral Shah
Keno Fischer
Rory Finnegan
Mike Innes
Glenn Moynihan

Dhairya Gandhi
Branwen Snelling
Gaurav Dhingra
Niklas Heim
Andrew Rosemberg
Niklas Schmitz
Chris Rackauckas
Andrei Zhabinski
Mathieu Besançon
Jeffrey Sarnoff
Anton Isopoussu
Antoine Levitt

Fernando Chorney
James Bradbury
Ben Cottier
Oliver Schulz
Alex Robson
cormullion
Roger Luo
Simon Etter
ho-oto
Wessel Bruinsma
Takafumi Arakaki
Carlo Lucibello
Eric Davies



Tour de ChainRules



Introduction

When to write rules

How to write rules

Basic example

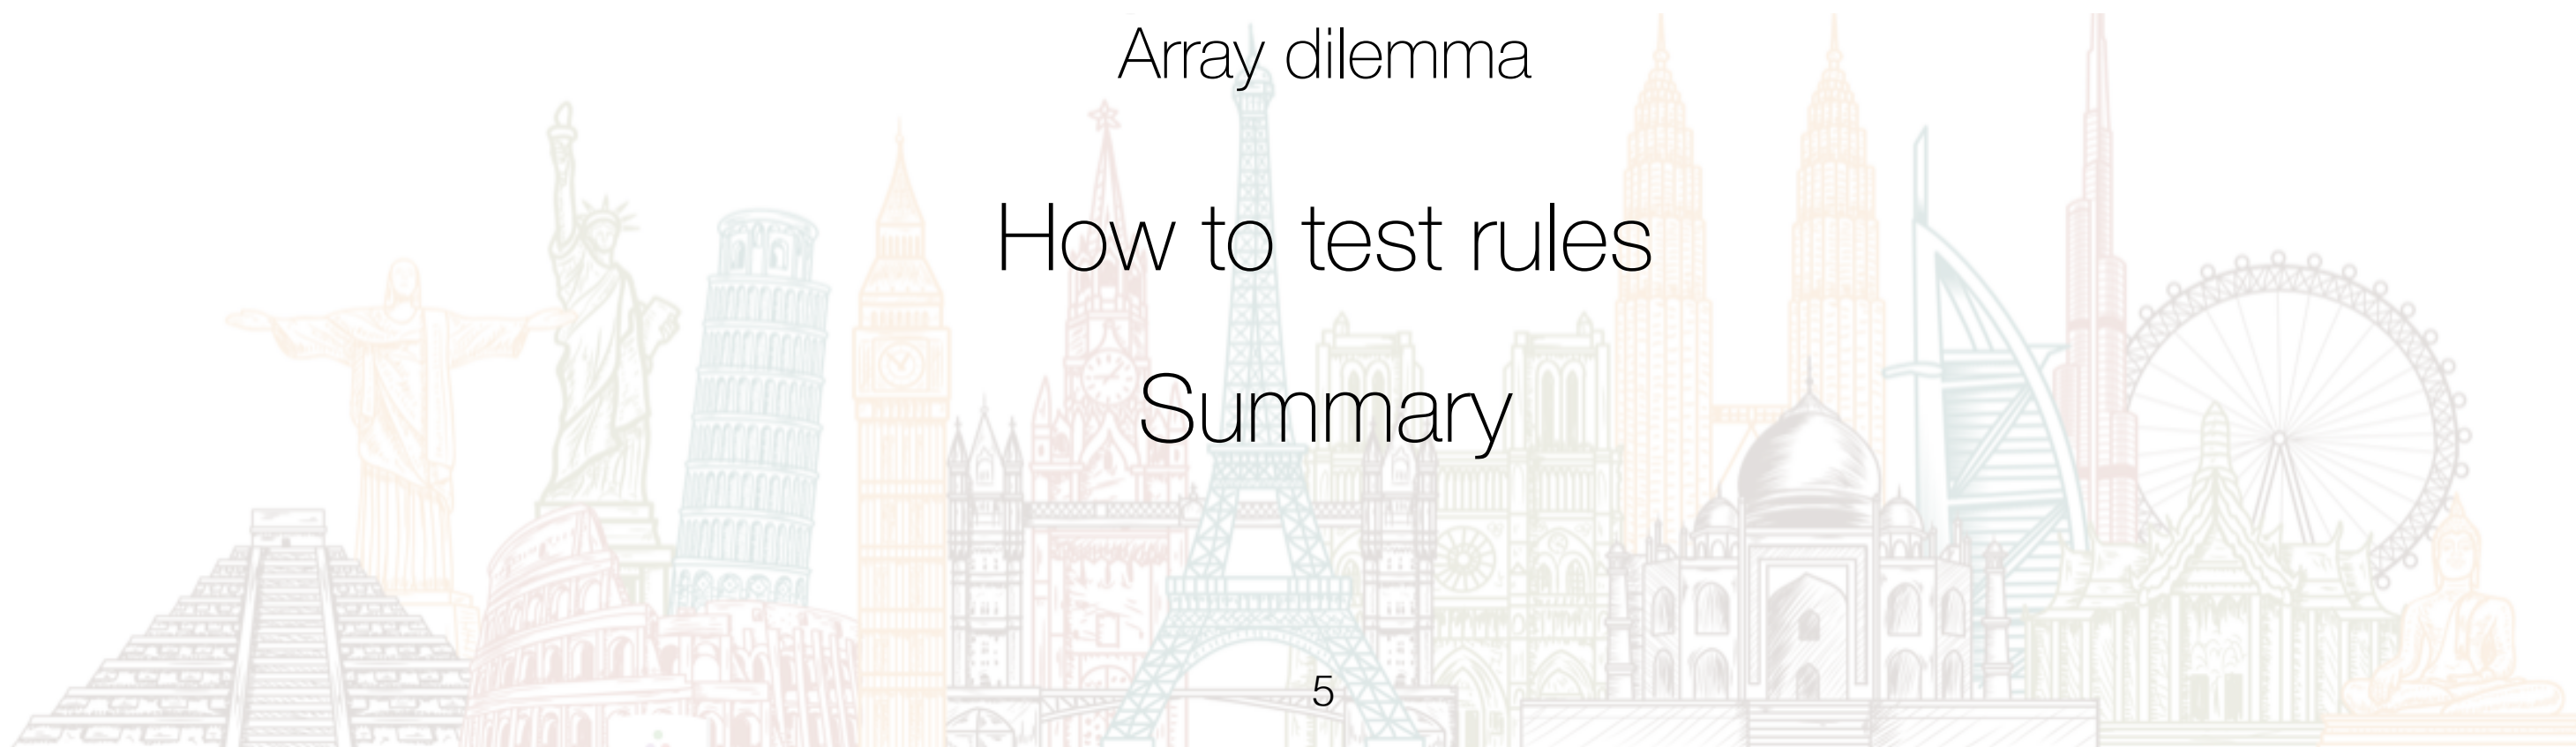
Gotchas and helpers

RuleConfig

Array dilemma

How to test rules


Summary



[Lyndon's talk last JuliaCon](#) is a great introduction


But in a nutshell:

- Automatic differentiation (AD) is useful
- We have many AD systems in Julia
- They all need rules that are efficient and correct




ChainRules.jl

(rules for Julia stdlibs)



ChainRulesCore.jl

(utilities to define rules)




ChainRulesTestUtils.jl

(utilities for testing rules)

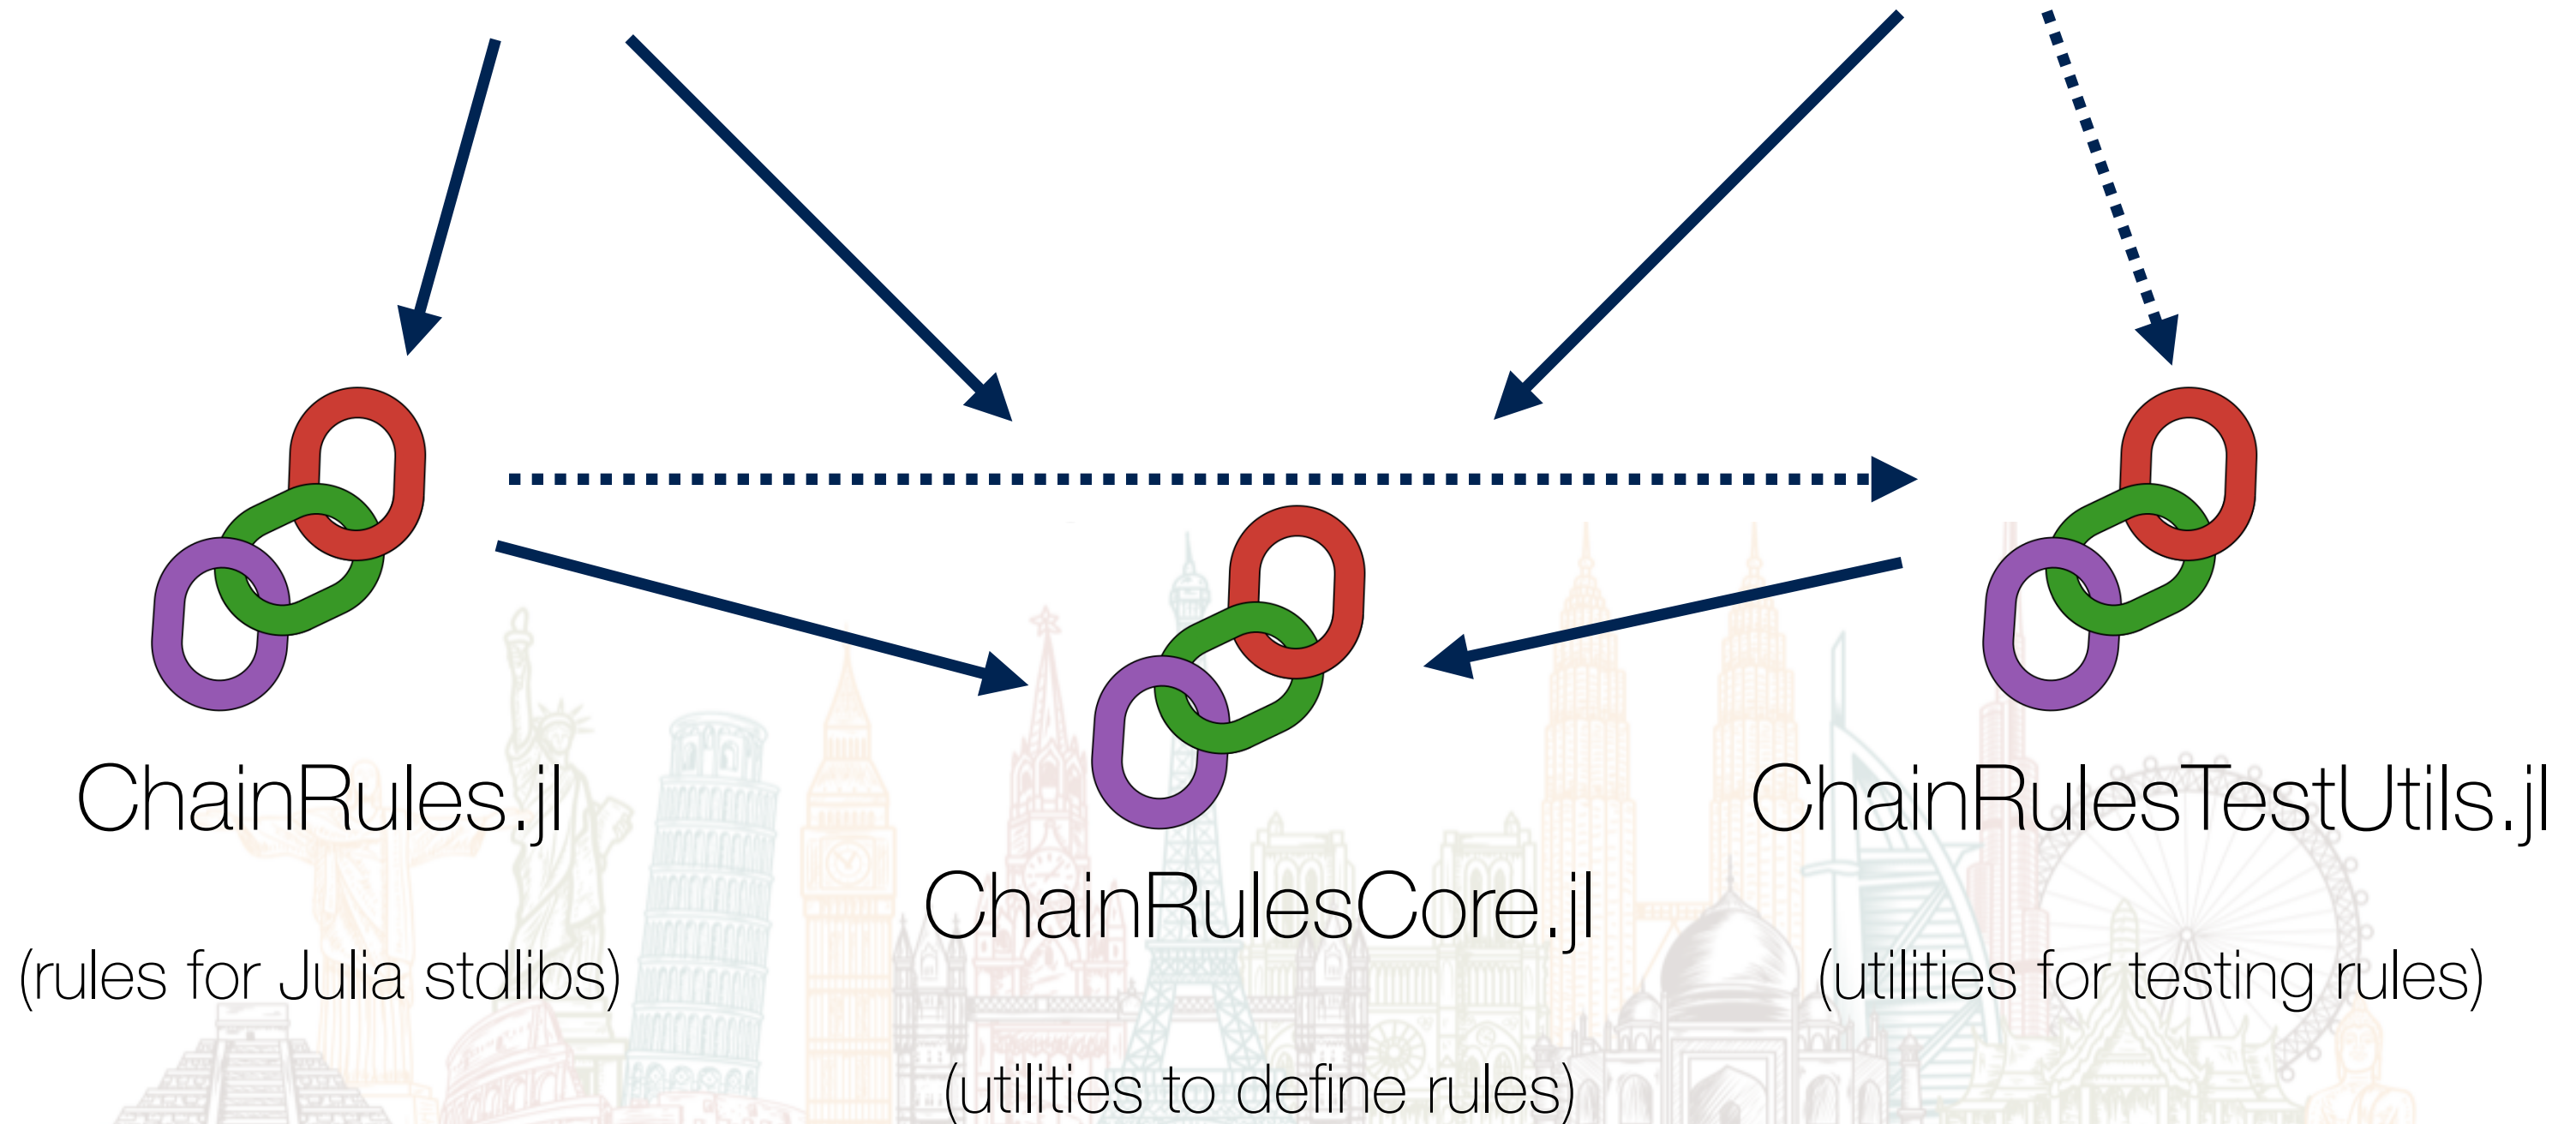
AD packages



other packages



SpecialFunctions.jl



Kinds of rules

frule

(Forwards mode AD)

```
function frule((f̄, args...), f, args...)  
    y = f(args)  
    ...  
    return y, ȳ  
end
```

rrule

(Backwards mode AD)

```
function rrule(f, args...)  
    y = f(args...)  
    function f_pullback(ȳ)  
        ...  
        return (f̄, āargs...)  
    end  
    return y, f_pullback  
end
```

$$\dot{x} = \frac{\partial x}{\partial(\text{something})}$$

$$\bar{x} = \frac{\partial(\text{something})}{\partial x}$$

Kinds of tangents

ZeroTangent() can be perturbed, but no change in output

NoTangent() can not be perturbed

Tangent{Foo}() tangent of **structs**, **Tuples**, **NamedTuples**, and **Dicts**

Thunk() delayed computation


InplaceableThunk() delayed computation, can accumulate gradients inplace



Tour de ChainRules



Introduction



When to write rules

How to write rules

Basic example

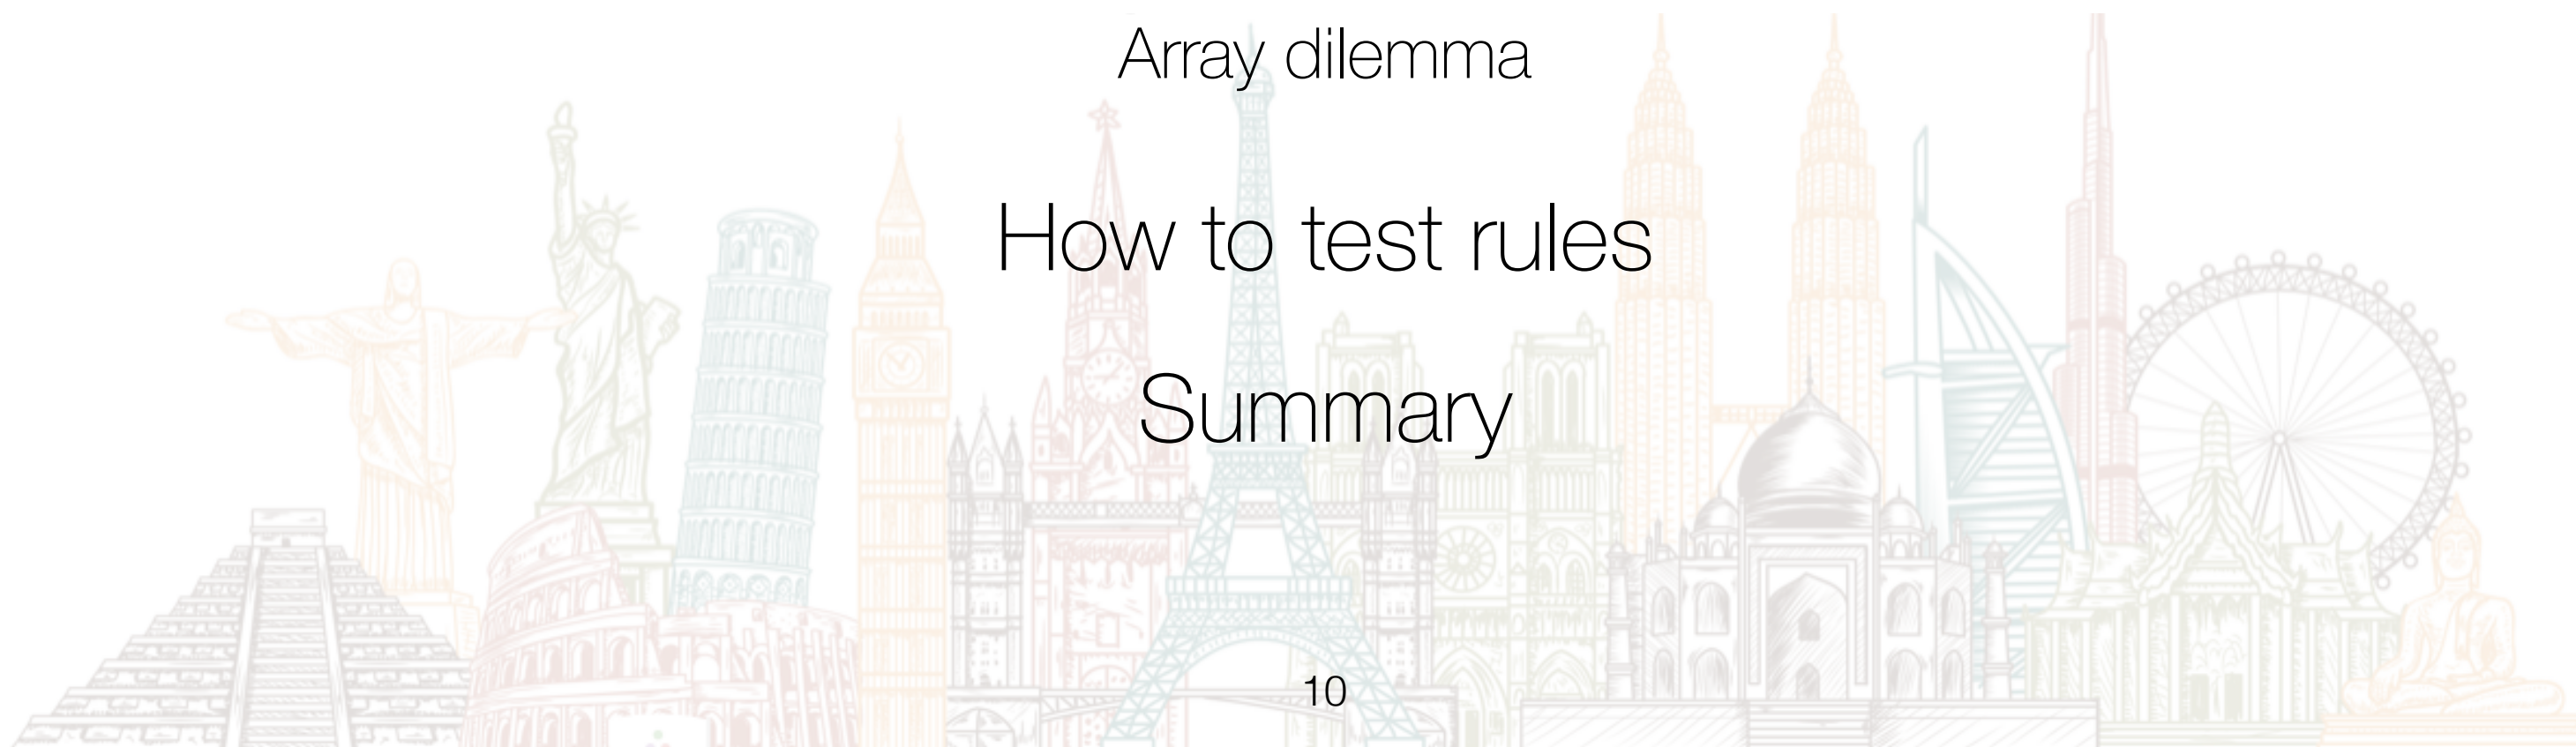
Gotchas and helpers

RuleConfig

Array dilemma

How to test rules

Summary



perfect AD system needs rules

to work at all: basic operations (+, *, ...)

to be efficient: e.g. AD through numerical integration, AD through DiffOpt

imperfect AD systems need more rules

to work at all: unsupported features (e.g. mutation)

to be efficient: e.g. differentiating through loops



When to write rules for Zygote

To work around unsupported features

```
function mutation(n::Integer)
    array = zeros(n)
    array[1] = 1
    return sum(array)
end
```

```
function exception_handling(x)
    try
        return x^2
    catch e
        println("can't square $x")
        throw(e)
    end
end
```

To improve efficiency

```
function no_inplace_accum(array)
    x = array[1]
    y = array[2]
    z = array[3]
    return x+y+z
end
```

```
function for_loops(array)
    s = 0
    for a in array
        s += a
    end
    return s
end
```


When to write rules for Zygote

To work around unsupported features

```
function mutation(n::Integer)
    array = zeros(n)
    array[1] = 1
    return sum(array)
end
```

```
function exception_handling(x)
    try
        return x^2
    catch e
        println("can't square $x")
        throw(e)
    end
end
```

To improve efficiency

```
function no_inplace_accum(array)
    x = array[1]
    y = array[2]
    z = array[3]
    return x+y+z
end
```

```
function for_loops(array)
    s = 0
    for a in array
        s += a
    end
    return s
end
```

When to write rules for Zygote

To work around unsupported features

```
function mutation(n::Integer)
    array = zeros(n)
    array[1] = 1
    return sum(array)
end
```

```
function exception_handling(x)
    try
        return x^2
    catch e
        println("can't square $x")
        throw(e)
    end
end
```

To improve efficiency

```
function no_inplace_accum(array)
    x = array[1]
    y = array[2]
    z = array[3]
    return x+y+z
end
```

```
function for_loops(array)
    s = 0
    for a in array
        s += a
    end
    return s
end
```





Tour de ChainRules



Introduction

When to write rules



How to write rules



Basic example


Gotchas and helpers

RuleConfig

Array dilemma

How to test rules

Summary



Primal program

just a matrix multiplication
(note that `Foo.c` is not involved)

```
struct Foo
  A::Matrix
  c::Float64
end
```

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```



Code to write the rrule

```
struct Foo
  A::Matrix
  c::Float64
end
```

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```

```
import ChainRulesCore: rrule
```

need to extend the function



Code to write the rrule

```
struct Foo
  A::Matrix
  c::Float64
end
```

Reminder of the rrule signature

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```

```
import ChainRulesCore: rrule

function rrule(::typeof(foo_mul), foo::Foo, B::AbstractArray)
  Y = foo_mul(foo, B)
  function foo_mul_pullback(Ȳ)
    f̄ = NoTangent()
    f̄foo = Tangent{Foo} (; A=Ȳ * B', c=ZeroTangent())
    B̄ = foo.A' * Ȳ
    return f̄, f̄foo, B̄
  end
  return Y, foo_mul_pullback
end
```

Code to write the rrule

```
struct Foo
  A::Matrix
  c::Float64
end
```

It is possible to change the primal computation
(might want for efficiency if work shared in primal and pullback)

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```

```
import ChainRulesCore: rrule

function rrule(::typeof(foo_mul), foo::Foo, B::AbstractArray)
  Y = foo_mul(foo, B)
  function foo_mul_pullback(Ȳ)
    f̄ = NoTangent()
    f̄foo = Tangent{Foo} (; A=Ȳ * B', c=ZeroTangent())
    B̄ = foo.A' * Ȳ
    return f̄, f̄foo, B̄
  end
  return Y, foo_mul_pullback
end
```


Code to write the rrule

```
struct Foo
  A::Matrix
  c::Float64
end
```

pullback signature:

primal output tangent -> primal input tangents

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```

```
import ChainRulesCore: rrule

function rrule(::typeof(foo_mul), foo::Foo, B::AbstractArray)
  Y = foo_mul(foo, B)
  function foo_mul_pullback(Ȳ)
    f̄ = NoTangent()
    f̄foo = Tangent{Foo} (; A=Ȳ * B', c=ZeroTangent())
    B̄ = foo.A' * Ȳ
    return f̄, f̄foo, B̄
  end
  return Y, foo_mul_pullback
end
```

Code to write the rrule

```
struct Foo
    A::Matrix
    c::Float64
end
```

`foo_mul` does not have fields, use **NoTangent()**

NB: functors (callable structs) can have fields

```
function foo_mul(foo::Foo, B::AbstractArray)
    return foo.A * B
end
```

```
import ChainRulesCore: rrule

function rrule(::typeof(foo_mul), foo::Foo, B::AbstractArray)
    Y = foo_mul(foo, B)
    function foo_mul_pullback(Ȳ)
        f̄ = NoTangent()
        f̄foo = Tangent{Foo} (; A=Ȳ * B', c=ZeroTangent())
        B̄ = foo.A' * Ȳ
        return f̄, f̄foo, B̄
    end
    return Y, foo_mul_pullback
end
```

Code to write the rrule

```
struct Foo
  A::Matrix
  c::Float64
end
```

tangent of **foo::Foo** is **Tangent{Foo} (; ...)**
tangent of the **c** field is **ZeroTangent()**

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```

```
import ChainRulesCore: rrule

function rrule(::typeof(foo_mul), foo::Foo, B::AbstractArray)
  Y = foo_mul(foo, B)
  function foo_mul_pullback(Ȳ)
    f̄ = NoTangent()
    f̄foo = Tangent{Foo} (; A=Ȳ * B', c=ZeroTangent())
    B̄ = foo.A' * Ȳ
    return f̄, f̄foo, B̄
  end
  return Y, foo_mul_pullback
end
```

Non-specified **Tangent** fields are implicitly **ZeroTangent()**

Code to write the rrule

```
struct Foo
  A::Matrix
  c::Float64
end
```

use **Thunk**s (to delay and potentially avoid computation)
or **InplaceableThunks** (to accumulate gradients inplace)

e.g. $\bar{B} = @thunk \text{foo}.A' * \bar{Y}$

```
function foo_mul(foo::Foo, B::AbstractArray)
  return foo.A * B
end
```

```
import ChainRulesCore: rrule

function rrule(::typeof(foo_mul), foo::Foo, B::AbstractArray)
  Y = foo_mul(foo, B)
  function foo_mul_pullback( $\bar{Y}$ )
     $\bar{f}$  = NoTangent()
     $\bar{\text{foo}}$  = Tangent{Foo} (; A= $\bar{Y} * B'$ , c=ZeroTangent())
     $\bar{B} = \text{foo}.A' * \bar{Y}$ 
    return  $\bar{f}$ ,  $\bar{\text{foo}}$ ,  $\bar{B}$ 
  end
  return Y, foo_mul_pullback
end
```




Tour de ChainRules



Introduction

When to write rules



How to write rules

Basic example



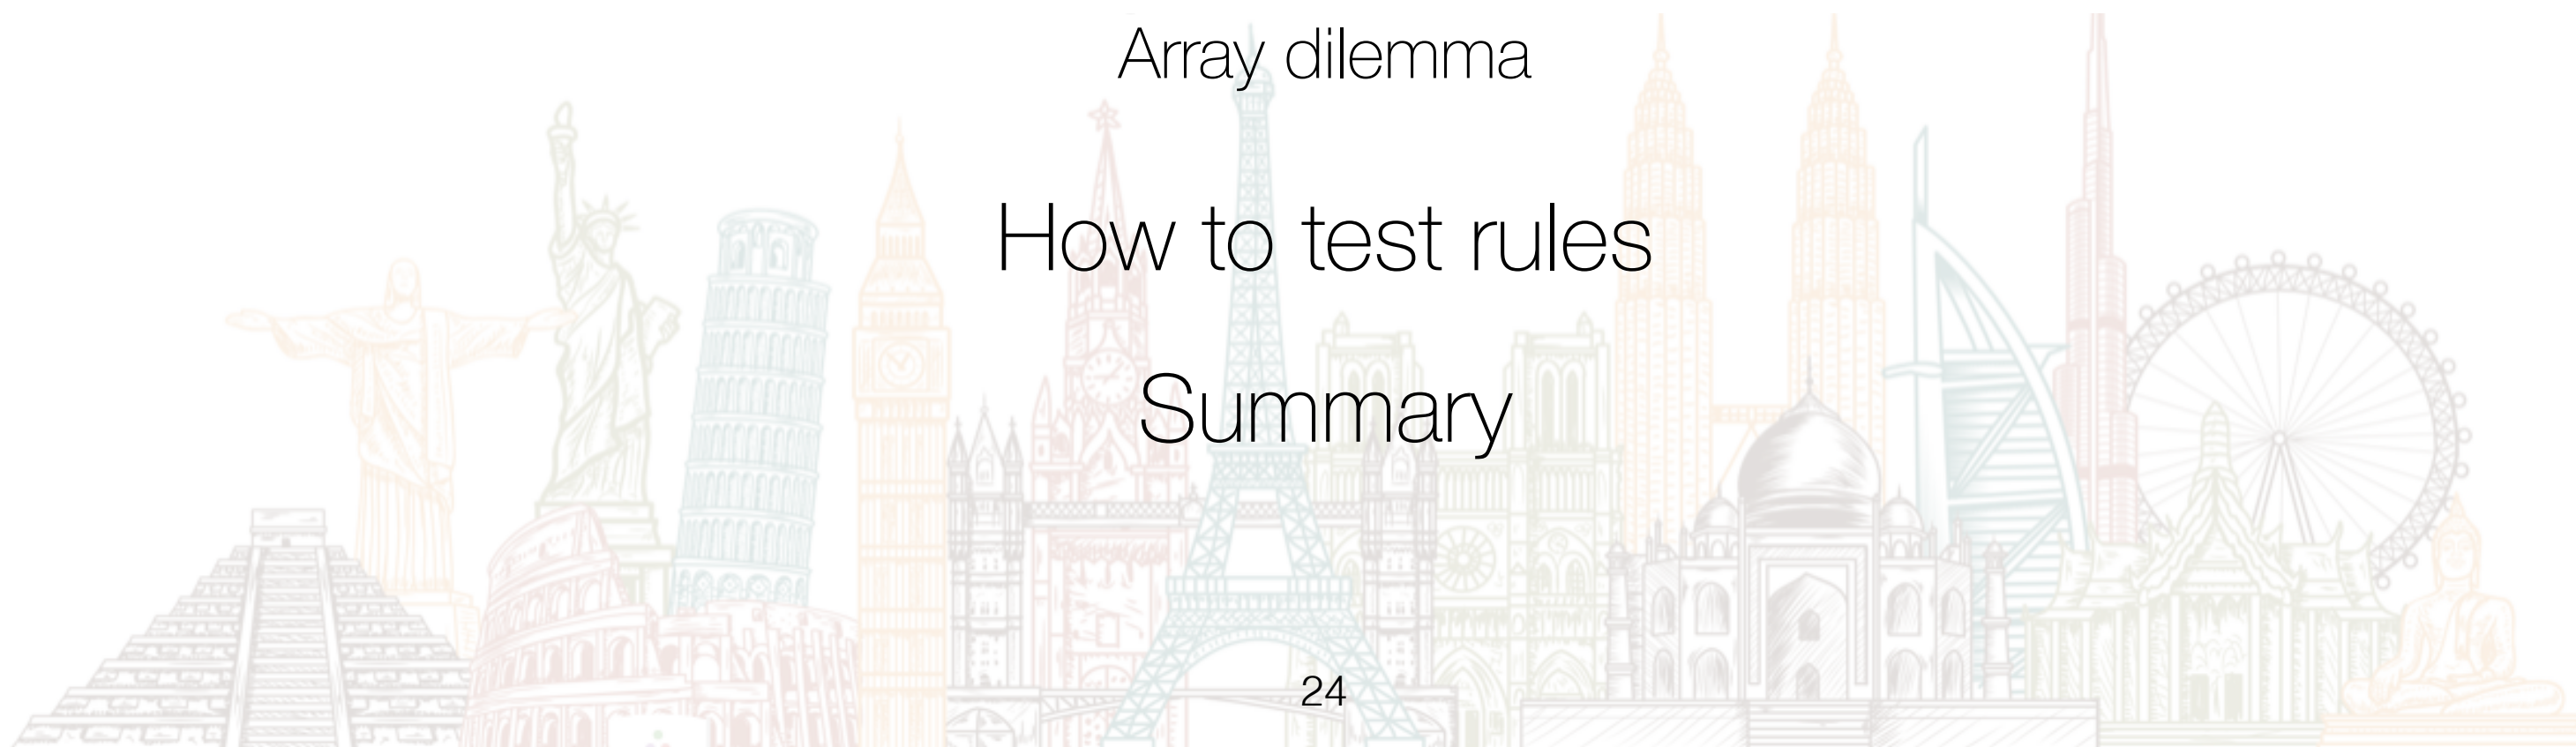
Gotchas and helpers

RuleConfig

Array dilemma

How to test rules

Summary



Struct gotchas

```
struct Foo
  A::Matrix
  c::Float64
end
```

function

```
func_foo(foo::Foo, B) = 77
```

```
rrule(::typeof(func_foo), f::Foo, B)
```

callable struct

```
(foo::Foo)(B) = 77
```

```
rrule(foo::Foo, B)
```

constructor

```
Foo(A) = Foo(A, 0.0)
```

```
rrule(::Type{Foo}, A)
```


note that

`typeof(Foo) == DataType`

so `::typeof(Foo)` defines an `rrule` for all constructors!

If a function has multiple inputs or outputs, and only some of the derivatives have been worked out analytically, **@not_implemented** macro can be used

```
 $\bar{B}$  = @not_implemented(  
    ""  
    The derivative w.r.t. B is not implemented:  
    https://github.com/User/Package.jl/issues/77  
    ""  
)
```



If all of the arguments of the function are non-perturbable,
@non_differentiable macro can be used
to automatically generate both the **frule** and the **rrule**

```
@non_differentiable joinpath(::AbstractString, ::AbstractString...)
```



If the function is defined on scalars*, i.e. `<:Number`,
`@scalar_rule` macro can be used
to automatically generate both the `frule` and the `rrule`

```
@scalar_rule sinh(x) cosh(x)
```






Tour de ChainRules



Introduction

When to write rules



How to write rules

Basic example

Gotchas and helpers




RuleConfig

Array dilemma

How to test rules

Summary



RuleConfig and calling back into AD

RuleConfig is a way to make rules conditionally defined, depending on the properties of the AD system.

This allows us to define a rule which calls forward mode AD inside the rule definition. This rule will only be used by AD systems which support forward mode.



Implemented with a trait-like system (not Holy Traits)

AD systems declare their properties

```
struct MyADRuleConfig <: RuleConfig{Union{Feature1, Feature2}} end
```

While rule authors can specify which properties are needed for the rule to be defined

```
# rrule that is only defined for ADs with `Feature1`  
rrule(::RuleConfig{>:Feature1, }, f, args...) = ...  
  
# frule that is only defined for ADs with both `Feature1` and `Feature2`  
frule(::RuleConfig{>:Union{Feature1, Feature2}}, f, args...) = ...
```


Calling back into AD

Complementary properties: **HasReverseMode**, **NoReverseMode**
and similarly for forwards mode

AD that **HasReverseMode** needs to define **rrule_via_ad** for its RuleConfig subtype:

```
struct MyReverseOnlyADRuleConfig <: RuleConfig{Union{HasReverseMode, NoForwardsMode}} end  
function ChainRulesCore.rrule_via_ad(::MyReverseOnlyADRuleConfig, f, args...)  
    ...  
    return y, pullback  
end
```



Rules for higher order functions


```
function rrule(  
  config::RuleConfig{>:HasForwardsMode},  
  ::typeof(map),  
  f::Function,  
  x::Array{<:Real}  
)  
  
  y_and_ẏ = map(x) do xi  
    frule_via_ad(config, (NoTangent(), one(xi)), f, xi)  
  end  
  y = first.(y_and_ẏ)  
  ẏ = last.(y_and_ẏ)  
  
  pullback_map(ȳ) = NoTangent(), NoTangent(), ȳ .* ẏ  
  return y, pullback_map  
end
```

Writing rules only for your own AD

Just dispatch on the RuleConfig (do not define a new feature)

```
struct MyADConfig <: RuleConfig{Union{Feature1, Feature2}} end
```

```
rrule(::MyADConfig, typeof(some_func), args...) = ...
```






Tour de ChainRules



Introduction

When to write rules



How to write rules

Basic example

Gotchas and helpers


RuleConfig



Array dilemma

How to test rules

Summary




Array dilemma

More formally: taking types that represent embedded subspaces seriously

⚠ This is complicated.

If you don't understand, it is my fault, not yours.

I will make sure to wake you up when we move on.



Primal computation

```
function sum_array(A::AbstractArray)
    s = 0
    for i in eachindex(A)
        s += A[i]
    end
    return s
end
```

```
function sum_array(A::Diagonal)
    return sum_array(diag(A))
end
```

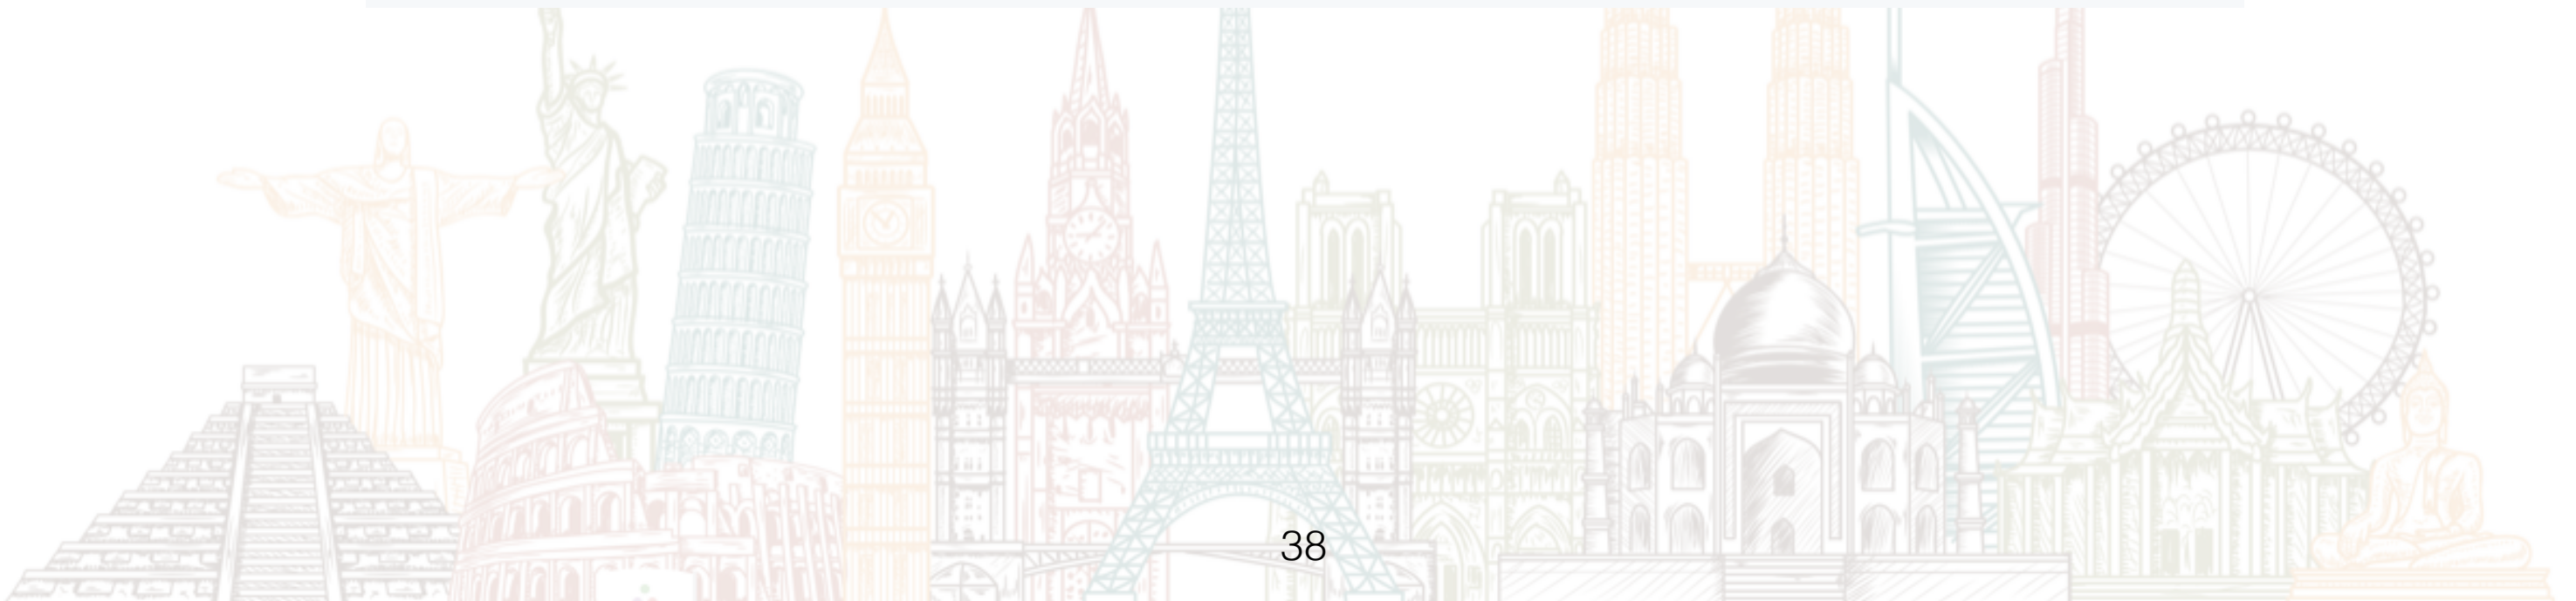
Summing an AbstractArray

Dispatch allows us to specialise

Pullback

For **A::Matrix**, which uses the fallback primal, the pullback is just a matrix of \bar{y} :

```
function rrule(::typeof(sum_array), A::AbstractArray)
    y = sum_array(A)
    sizeA = size(A)
    function sum_array_pullback( $\bar{y}$ )
        return NoTangent(), fill( $\bar{y}$ , sizeA)
    end
    return y, sum_array_pullback
end
```




Pullback

What do we want for **A::Diagonal**, which uses the specialised primal?

Option 0:

We can define a custom rule.

But this approach means we have to write very many rules.
(and add new rules as new types appear)



Pullback


What do we want for **A::Diagonal**, which uses the specialised primal, if we have not defined a custom **rrule**?

Option 1:

Fallback to abstractly typed **rrule** with **A::AbstractArray**.

Option 2:

Let AD have a go at differentiating the specialised primal computation.



Pullback

What do we want for **A::Diagonal**, which uses the specialised primal, if we have not defined a custom **rrule**?

Option 1:

Fallback to abstractly typed **rrule** with **A::AbstractArray**.

⚠ can return the wrong answer: e.g. a dense matrix tangent for **A::Diagonal**

⚠ can be slower than Option 2

Option 2:

Let AD have a go at differentiating the specialised primal computation.

⚠ can error if the specialised primal computation uses unsupported features (e.g. mutation)

⚠ can be slower Option 1 (e.g. for loops are slow)

Pullback

What do we want for **A::Diagonal**, which uses the specialised primal, if we have not defined a custom **rrule**?

Option 1:

Fallback to abstractly typed **rrule** with **A::AbstractArray**.

⚠ can return the wrong answer: e.g. a dense matrix tangent for **A::Diagonal**

⚠ can be slower than Option 2

Option 2:

~~Let AD have a go at differentiating the specialised primal computation.~~

~~⚠ can error if the specialised primal computation uses unsupported features (e.g. mutation)~~

~~⚠ can be slower Option 1 (e.g. for loops are slow)~~

Pullback

What do we want for **A::Diagonal**, which uses the specialised primal, if we have not defined a custom **rrule**?

Option 1:

Fallback to abstractly typed **rrule** with **A::AbstractArray**.

⚠ can return the wrong answer: e.g. a dense matrix tangent for **A::Diagonal**

⚠ can be slower than Option 2

To fix these problems:
ProjectTo and **@opt_out**

⚠ can return the wrong answer: e.g. a dense matrix tangent for **A::Diagonal**

Solve by making sure the tangent remains in the subspace of the primal:

create an object that knows how to project on the right tangent type
(i.e. knows the type, but also the size etc.)

project_A = ProjectTo(A)

and then project the tangent

project_A(tangent)

```
function rrule(::typeof(sum_array), A::Matrix)
    y = sum_array(A)
    sizeA = size(A)
    project_A = ProjectTo(A)
    function sum_array_pullback(ȳ)
        return NoTangent(), project_A(fill(ȳ, sizeA))
    end
    return y, sum_array_pullback
end
```


⚠ can be slower than AD through specialised primal computation

Solve by providing a way to opt out of using the fallback rule:

```
@opt_out rrule(::typeof(sum_array), A::Diagonal)
```

for a particular function signature.

This makes it easy to check whether AD or fallback rule is faster.





Tour de ChainRules



Introduction

When to write rules


How to write rules

Basic example

Gotchas and helpers


RuleConfig

Array dilemma



How to test rules

Summary



Testing rules

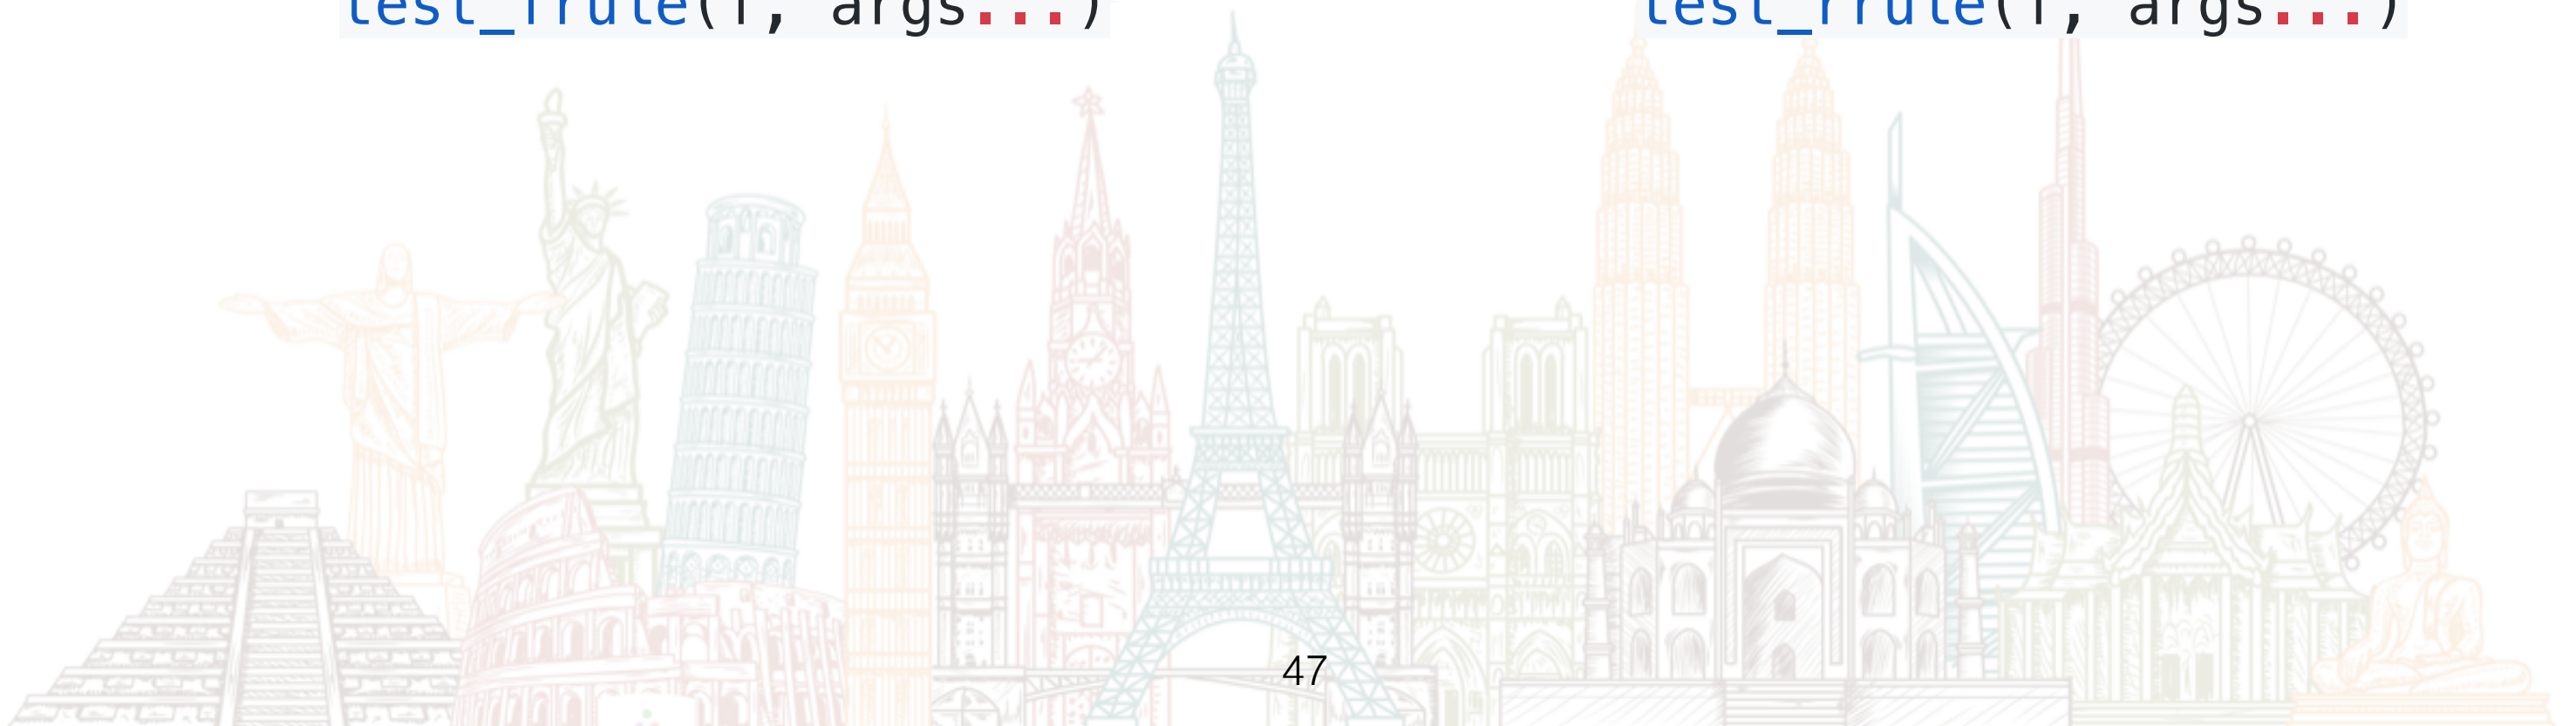
Need to test rules with finite differencing methods.

frule

```
frule((f, args...), f, args...)
test_frule(f, args...)
```

rrule

```
rrule(f, args...)
test_rrule(f, args...)
```



Testing rules

Powered by automatic tangent generation

```
test_rrule(f, x)
```

Can* specify the tangent explicitly: `primal \vdash tangent`


```
test_rrule(f  $\vdash$   $\bar{f}$ , x  $\vdash$   $\bar{x}$ ; output_tangent= $\bar{y}$ )
```

*sometimes *have to*

Testing AD gradients

Specify the rrule-like function

```
test_rrule(f, args...; rrule_f=rrule_via_ad)
```





Tour de ChainRules



Introduction

When to write rules

How to write rules


Basic example

Gotchas and helpers


RuleConfig

Array dilemma

How to test rules



Summary



ChainRules 1.0 is out!



new features

used by AD systems

many rules

calling back into AD

100s of rules in ChainRules.jl

ProjectTo

Diffractor

AbstractGPs.jl

@opt_out

Nabla

BlockDiagonals.jl

pullbacks handle Thunks

ReversePropagation

DiffEqBase.jl

testing rules is easier

Yötä

Hankel.jl

can test AD gradients

Zygote


PDMatsExtras.jl

@non_differentiable


SpecialFunctions.jl

@not_implemented

...



ChainRules 1.x/2.0 wish list



Higher order rules (2nd derivative)

Rules for mutating functions

Better solution for the array dilemma

Better solution for inplace accumulation

Wirtinger derivatives for complex numbers

Rules for jacobians

Ability to get a basis for anything

Tooling to support systems that can't augment the primal

Summary

 ChainRules project is at 1.0 

Integrated in many AD systems

100s of correct and efficient rules

Utilities to easily write more rules

Utilities to test rules (and AD gradients) using finite differences